

Using AWS CodePipeline for secure CloudFormations deployments

AWS offers a rich set of developer tools to host code, build, and deploy your application and/or infrastructure to AWS. These include [AWS CodePipeline](#), for continuous integration and continuous deployment orchestration; [AWS CodeCommit](#), a fully-managed source control service; [AWS CodeBuild](#), a fully-managed continuous integration service that compiles source code, runs tests, and produces software packages; and [AWS CodeDeploy](#), a fully-managed deployment service that automates software deployments to a variety of compute services. These services can be used together to build powerful CI/CD pipelines on AWS.

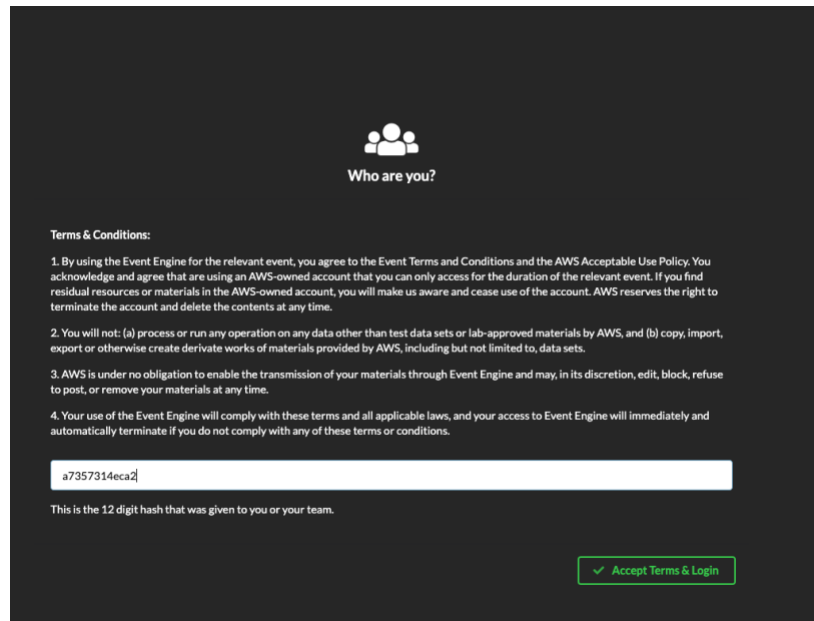
In addition, AWS developer tools allow you to leverage open source software (OSS) to build CI/CD pipelines on AWS, so you can continue to use the tools you are familiar with. For instance, you can use GitHub as the source code provider for AWS CodePipeline or run CodeBuild jobs orchestrated by Jenkins. In fact, a variety of OSS and third-party tools can be easily integrated with AWS CodeBuild, allowing you to extend the capabilities of your pipeline. As an extra benefit, you are not required to manage build servers since CodeBuild is a fully-managed service.

In this lab, we will show you how to build a serverless infrastructure deployment pipeline (i.e., no need for you to manage build servers) on AWS using AWS developer tools in conjunction with popular open source tools such as CFN-Nag, CFN-Python-Lint, and Stacker. The pipeline will run automated validation checks against CloudFormation templates and deploy the corresponding CloudFormation stacks if the templates are valid.

The companion source code for this Lab can be found in the [GitHub repo `aws-codepipeline-secure-cfn-lab`](#).

Lets's get started!

1. Go to <https://dashboard.eventengine.run/login> and enter the hash you've received and Login



The login page has a dark background. At the top center is an icon of three people and the text "Who are you?". Below this is a section titled "Terms & Conditions:" followed by four numbered terms. A text input field contains the hash "a7357314eca2". Below the input field is a line of text: "This is the 12 digit hash that was given to you or your team." At the bottom right is a green button with a checkmark icon and the text "Accept Terms & Login".

Who are you?

Terms & Conditions:

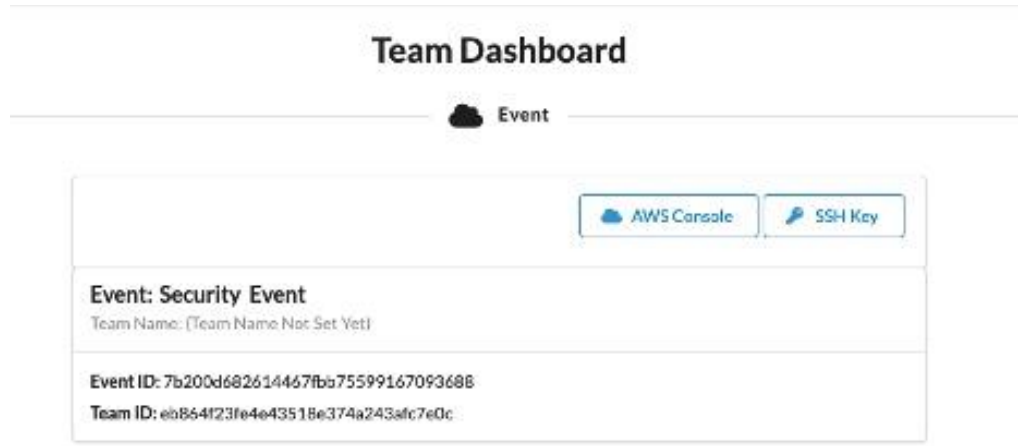
1. By using the Event Engine for the relevant event, you agree to the Event Terms and Conditions and the AWS Acceptable Use Policy. You acknowledge and agree that are using an AWS-owned account that you can only access for the duration of the relevant event. If you find residual resources or materials in the AWS-owned account, you will make us aware and cease use of the account. AWS reserves the right to terminate the account and delete the contents at any time.
2. You will not: (a) process or run any operation on any data other than test data sets or lab-approved materials by AWS, and (b) copy, import, export or otherwise create derivate works of materials provided by AWS, including but not limited to, data sets.
3. AWS is under no obligation to enable the transmission of your materials through Event Engine and may, in its discretion, edit, block, refuse to post, or remove your materials at any time.
4. Your use of the Event Engine will comply with these terms and all applicable laws, and your access to Event Engine will immediately and automatically terminate if you do not comply with any of these terms or conditions.

a7357314eca2

This is the 12 digit hash that was given to you or your team.

✓ Accept Terms & Login

You now have access to your Team dashboard. It will provide you with a link and credentials to the AWS console



The dashboard has a light gray header with the title "Team Dashboard" and a sub-header "Event" with a cloud icon. Below the header is a white box containing two buttons: "AWS Console" with a cloud icon and "SSH Key" with a key icon. Below the buttons is a section titled "Event: Security Event" with the text "Team Name: (Team Name Not Set Yet)". Below this is a section with "Event ID: 7b200d682614467fb575599167093688" and "Team ID: eb864f23fe4e43518e374a243afc7e0c".

Team Dashboard

Event

AWS Console SSH Key

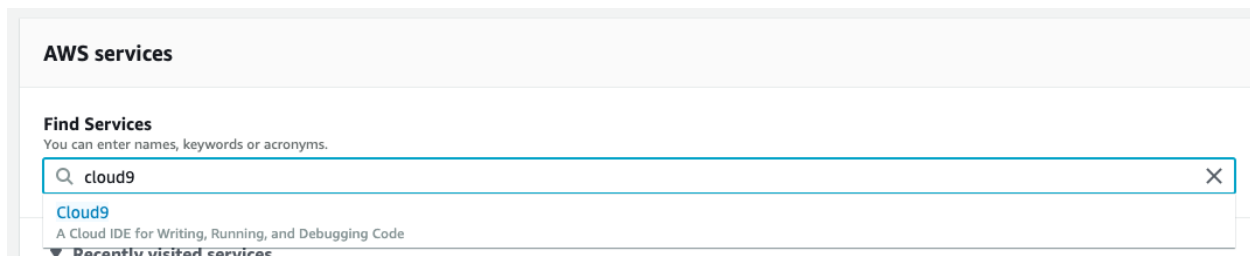
Event: Security Event

Team Name: (Team Name Not Set Yet)


Event ID: 7b200d682614467fb575599167093688

Team ID: eb864f23fe4e43518e374a243afc7e0c

2. Next you will need a development environment. We will use AWS Cloud9 for that. Enter Cloud9 into the search bar and press enter.

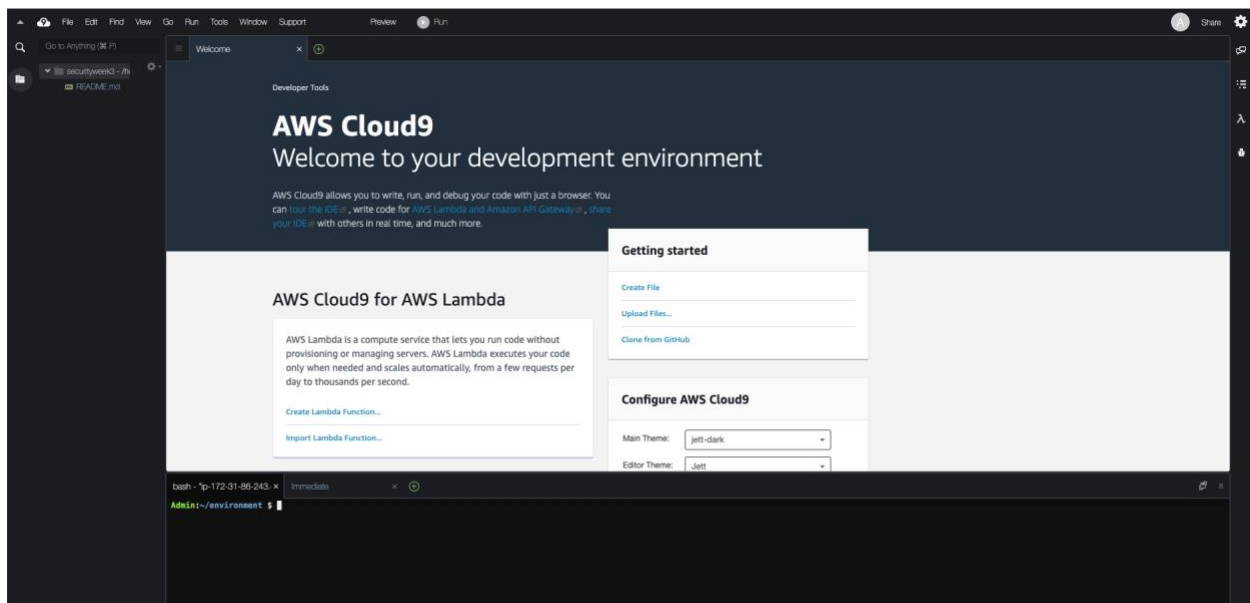


This will lead you to the Developer Tools page where you can click Create environment button.



Enter a name for your environment and click Next step on the bottom right.

For configuration settings leave the default, choose **Instance Type** *t3.small* and for **Platform** *Amazon Linux 2*. In terms of **cost saving settings** keep the default after **30 minutes**. Then create the environment. You will see the following screen when the environment is up:



Your environment is now ready.

In the next few steps we will install various tooling so that you can check your templates locally on your Cloud9 environment.

Open source tools

First, let's briefly review the open source tools that will be later used to build the infrastructure pipeline discussed in this lab.

CFN-Nag

[CFN-Nag](#) is a popular open source tool developed by [Stelligent](#) and provided to the open source community to help pinpoint security problems early on in an [AWS CloudFormation](#) template. CFN-Nag looks for patterns in AWS CloudFormation templates that may indicate insecure infrastructure, for example:

- IAM rules that are too permissive (wildcards)
- Security group rules that are too permissive (wildcards)
- Access logs that aren't enabled
- Encryption that isn't enabled

To install CFN-Nag from your workstation, you'll need Ruby v.2.5 or greater installed. Assuming that you have met this prerequisite, simply run:

```
Admin:~/environment $ gem install cfn-nag
...
Done installing documentation for kwalify, cfn-model, multi_json, little-plugger, logging, netaddr, optimist, jmespath,
aws-partitions, aws-eventstream, aws-sigv4, aws-sdk-core, aws-sdk-kms, aws-sdk-s3, lightly, cfn-nag after 8 seconds
16 gems installed
```

Later in the lab when you have a CloudFormations template you can run CFN-Nag from the command line as such:

```
Admin:~/environment $ cfn_nag_scan --input-path <path to cloudformations json>
```

CFN-Lint

[CFN-Python-Lint](#) was released by AWS to the open source community as a CloudFormation linter. It validates AWS CloudFormation templates (YAML and JSON) against the AWS CloudFormation spec and does additional checks including ensuring valid values for resource properties, and best practices – for instance, if an input parameter is defined in a template but never referenced, the linter will raise a warning.

Install CFN-Python-Lint in your Cloud9 environment you can use pip3:

```
Admin:~/environment $ pip3 install cfn-lint
...
Successfully installed attrs-20.2.0 aws-sam-translator-1.27.0 boto3-1.15.17 botocore-1.18.18 cfn-lint-0.38.0
decorator-4.4.2 jmespath-0.10.0 jsonpatch-1.26 jsonpointer-2.0 jsonschema-3.2.0 junit-xml-1.9 networkx-2.5
pyparsing-2.4.7 python-dateutil-2.8.1 pyyaml-5.3.1 s3transfer-0.3.3 six-1.15.0 urllib3-1.25.10
```

To invoke CFN-Python-Lint from the command line later in the lab you can run the following command:

```
Admin:~/environment $ cfn-lint <path to yaml template>
```

CFN-Guard

[AWS CloudFormation-Guard](#) was released by AWS to the open source community as a CLI tool that checks AWS CloudFormation templates for policy and compliance using a simple, policy-as-code, declarative syntax. CFN-guard can replace CFN-Nag. Compared to CFN-nag, cnf-guard does not require you to learn ruby or implement classes to check the things you want to validate. In this lab we will run both so you can get a comparison. CloudFormation-guard also allows you to autogenerate rules based of a CloudFormation's template.

For this lab, we have created a cfn-guard rpm which you can install directly from the Github repository.

```
Admin:~/environment $ sudo yum install -y cargo rust && sudo rpm -ivh https://bit.ly/31CMLw3
```

To invoke cnf-guard from the command line later in the lab you can run the following command:

```
Admin:~/environment $ cfn-guard check --rule_set test.rules --template yourtemplate.yaml
```

Stacker

[Stacker](#) is an open source tool and library created by the [Remind engineering team](#) and released to the open source community. It is used to orchestrate and manage CloudFormation stacks across one or more accounts. Stacker can manage stack dependencies automatically and allow the output parameters of a given stack to be passed as input to other stacks (even for stacks deployed in different accounts!) making sure the stacks are created in the right order. It can also parallelize the deployment of non-dependent stacks reducing deployment time significantly.

Install [Stacker](#) in your Cloud9 environment by using pip3:

```
Admin:~/environment $ pip3 install stacker
```

```
...
```

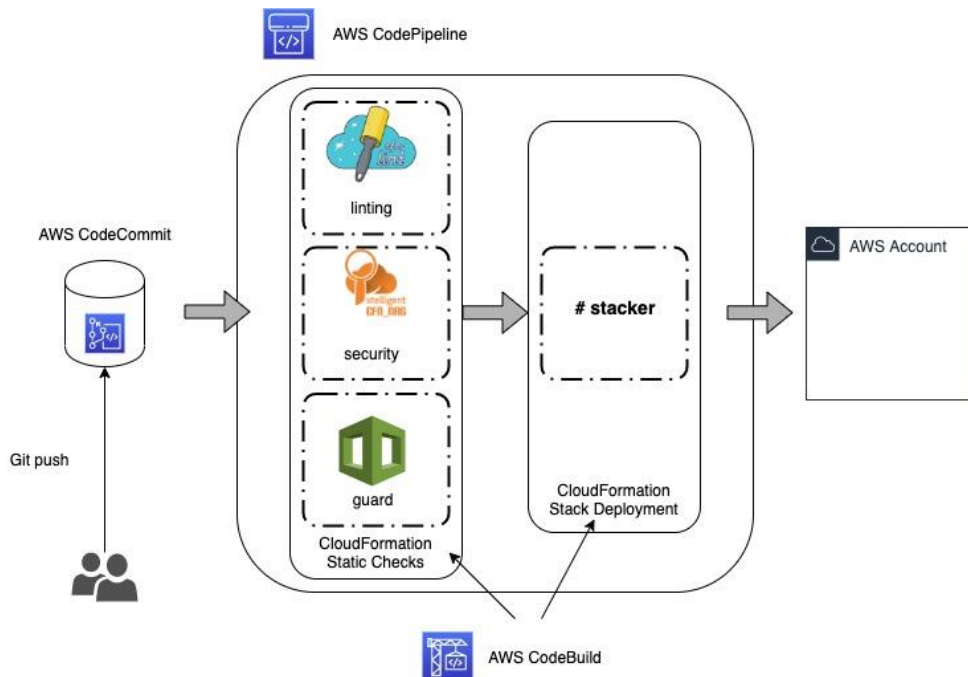
```
Successfully installed Click-7.1.2 MarkupSafe-1.1.1 awacs-1.0.0 cfn-flip-1.2.3 ecdsa-0.16.0 formic2-1.0.3 future-0.18.2 gitdb2-2.0.6 gitpython-2.1.15 jinja2-2.11.2 more-itertools-5.0.0 pyasn1-0.4.8 python-jose-3.1.0 rsa-4.5 schematics-2.0.1 smmap-3.0.4 smmap2-3.0.1 stacker-1.7.1 troposphere-2.6.3
```

To run this command later in the lab, use the following command:

```
Admin:~/environment $ stacker build <path to cloudformation yaml or json template>
```

Stacker will either launch or update your AWS CloudFormation stacks based on their configurations. Stacker can detect if the templates or their parameters have changed. If no changes are detected to a given stack, Stacker will not update that particular stack.

Overview of the secure infrastructure pipeline



The figure above shows the infrastructure deployment pipeline we will build in this lab. It uses AWS services such as AWS CodePipeline, AWS CodeCommit, and AWS CodeBuild in conjunction with open source software such as CFN-Nag, CFN-Python-Lint, CFN-Guard and Stacker.

The pipeline works as follows: a developer creates CloudFormation templates to provision infrastructure resources on AWS (e.g., a VPC with subnets, route tables, security groups, etc). When ready, the developer pushes the templates to an AWS CodeCommit repository which will then trigger the infrastructure pipeline. OSS tools CFN-Nag, CFN -Lint and CFN-Guard will run in parallel within a CodeBuild environment to validate the templates. CFN-Nag will check the templates against well-known security vulnerabilities while CFN-Python-Lint will enforce best practices and verify compliance with the CloudFormation specification. CFN-Guard will make sure that the templates are up to your defined standard/guardrails. If the templates do not pass the verification tests, the pipeline will stop in a failure state to prevent deployment issues. If the templates pass verification, Stacker will create or update CloudFormation stacks for each template. If Stacker fails, the pipeline will stop in a failure state to indicate a deployment issue. Otherwise, the stacks will be created or updated to provision the desired AWS resources.

As you might have noticed, the infrastructure pipeline described is generic and can be used to deploy arbitrary resources on AWS across a number of AWS accounts, while enforcing verification

rules against the submitted templates. For simplicity, we use Stacker to deploy to a single AWS account in this lab, but the tool really shines in the context of large cross-account deployments.

What will be deployed using the pipeline?

In this lab, we will use the infrastructure deployment pipeline to deploy three simple stacks as follows:

Stack #1: creates a private S3 bucket.

Stack #2: creates an EC2 IAM role that allows putting objects in the S3 bucket created in stack #1.

Stack #3: creates an EC2 instance and EC2 profile that assumes the IAM role created in stack #2 and uploads a simple file to the S3 bucket created in stack #1. Stack #3 depends on stacks #1 and #2.

Below is a snippet of the Stacker configuration file we use to create the three stacks we just mentioned:

```
#-----#
# Global variables
#-----#
namespace: cfnssec
stacker_execution_profile: &stacker_execution_profile stacker_execution
stacker_bucket: "" # not using S3 buckets to store CloudFormation templates

# any unique S3 bucket - suggestion: append your account id and region to the name
s3_bucket_name: &s3_bucket_name cfnssecsamples3bucket182743192031
# cheapest EC2 instance for testing purposes
ec2_instance_type: &ec2_instance_type t2.micro

#-----#
# Stack Definitions (https://stacker.readthedocs.io/en/latest/config.html)
#-----#
stacks:
- name: sample-s3-bucket
  profile: *stacker_execution_profile
  template_path: templates/s3-bucket-template.yaml
  variables:
    S3BucketName: *s3_bucket_name

- name: sample-ec2-iam-role
  profile: *stacker_execution_profile
  template_path: templates/ec2-role-template.yaml
  variables:
    S3BucketName: ${output sample-s3-bucket::S3BucketName}

- name: sample-ec2-instance
  profile: *stacker_execution_profile
  template_path: templates/ec2-template.yaml
  variables:
    InstanceType: *ec2_instance_type
    S3BucketName: ${output sample-s3-bucket::S3BucketName}
```

```
EC2IamRoleName: ${output sample-ec2-iam-role::EC2RoleName}
```

At the top of the file we define variables. The namespace variable is reserved for Stacker and used to prefix stack names. The other variables defined use [YAML anchors](#) that can be referenced in other parts of the template. For instance, the variable `ec2_instance_type` creates an anchor `ec2_instance_type` that refers to value “t2.micro”. This anchor value is later assigned to input variable `InstanceType` for stack `sample-ec2-instance` (see “`InstanceType: *ec2_instance_type`” above in the snippet of the Stacker configuration file).

The **stacks session** defines our three stacks. For each stack, we indicate the name of the stack, the AWS profile (credentials) that Stacker should use to deploy the stack, the local path to the CloudFormation template, and the input parameters to the stack. Notice that we use `${output ...}` to refer to stack output parameters and pass those as input parameters to other stacks. For instance, the `S3BucketName` output parameter from stack `sample-s3-bucket` is used as an input parameter by the other two stacks.

We start by provisioning the infrastructure pipeline that we introduced earlier.

3. First we need to get the source code. Clone the Git repository as shown below.

```
Admin:~/environment $ git clone https://github.com/ldomb/aws-codepipeline-secure-cfn-lab.git sample-cfnsec-pipeline
...
Cloning into 'sample-cfnsec-pipeline'...
remote: Enumerating objects: 32, done.
remote: Counting objects: 100% (32/32), done.
remote: Compressing objects: 100% (28/28), done.
remote: Total 32 (delta 7), reused 25 (delta 3), pack-reused 0
Unpacking objects: 100% (32/32), done.
```

Change into the directory you just cloned `sample-cfnsec-pipeline` directory and list the content:

```
Admin:~/environment $ cd sample-cfnsec-pipeline && ls -l
...
total 16
drwxrwxr-x 2 ec2-user ec2-user  24 Oct 17 11:45 cfnguard
-rw-rw-r-- 1 ec2-user ec2-user 309 Oct 17 11:45 CODE_OF_CONDUCT.md
drwxrwxr-x 2 ec2-user ec2-user   84 Oct 17 11:45 codepipeline
-rw-rw-r-- 1 ec2-user ec2-user 3306 Oct 17 11:45 CONTRIBUTING.md
-rw-rw-r-- 1 ec2-user ec2-user  932 Oct 17 11:45 LICENSE
-rw-rw-r-- 1 ec2-user ec2-user 1278 Oct 17 11:45 README.md
drwxrwxr-x 2 ec2-user ec2-user   96 Oct 17 11:45 rpms
drwxrwxr-x 2 ec2-user ec2-user   79 Oct 17 11:45 stacker
drwxrwxr-x 2 ec2-user ec2-user   92 Oct 17 11:45 templates
```

You should then see the following file structure:

- **codepipeline/**: CloudFormation templates to deploy the CI/CD pipeline including AWS CodePipeline, AWS CodeCommit, AWS CodeBuild and OSS CFN-Nag, CFN-Lint, and Stacker)
- **stacker**: Configuration files required by Stacker to perform CloudFormation stack deployments
- **stacker/buildspec.yaml**: CodeBuild buildspec that will install and invoke Stacker for CFN provisioning
- **stacker/stacker-config.yaml**: Stacker config file containing stack descriptions and input parameters
- **stacker/stacker-profiles**: AWS profiles file Stacker will use to deploy the various stacks
- **templates/***: the sample templates for the the stacks mentioned above that will be validated and deployed by the pipeline
- **rpms/**: This is the rpm to install cnf-guard
- **cfnguard**: here you will find a sample rules template

4. Create a local git repo for the pipeline artifacts

Now that you have the original source code, you'll create a local Git repository and copy the entire content of sample-cfnsec-pipeline over.

Assuming you have cloned the provided Git repository into folder **/home/ec2-user/environment/sample-cfnsec-pipeline/** you can make a directory named **/home/ec2-user/environment/sample-pipeline-artifacts/**. Then copy the entire directory sample-cfnsec-pipeline to this directory. The contents of this directory will be pushed to the repository that we will create shortly into AWS CodeCommit.

```
Admin:~/environment $ mkdir sample-pipeline-artifacts
Admin:~/environment $ cd sample-pipeline-artifacts
Admin:~/environment/sample-pipeline-artifacts $ git init
Admin:~/environment/sample-pipeline-artifacts $ cp -r ../sample-cfnsec-pipeline/* .
Admin:~/environment/sample-pipeline-artifacts $ git add .
```

We now added all the files to the local staging. We can check that with the following command

```
Admin:~/environment/sample-pipeline-artifacts $ git status
...
git status
On branch master

No commits yet

Changes to be committed:
  (use "git rm --cached <file>..." to unstage)
```

```
new file: CODE_OF_CONDUCT.md
new file: CONTRIBUTING.md
new file: LICENSE
new file: README.md
new file: cfnguard/test.rules
new file: codepipeline/codepipeline-template.yaml
new file: codepipeline/stacker-execution-role-template.yaml
new file: rpms/cfn-guard-1.0-0.amzn2.x86_64.rpm
new file: rpms/cfn-guard-debuginfo-1.0-0.amzn2.x86_64.rpm
new file: stacker/buildspec.yaml
new file: stacker/stacker-config.yaml
new file: stacker/stacker-profiles
new file: templates/ec2-role-template.yaml
new file: templates/ec2-template.yaml
new file: templates/s3-bucket-template.yaml
```

Before we commit anything, we need to setup our environment

```
Admin:~/environment/sample-pipeline-artifacts $ git config --global Firstname.Lastname "John Doe"
Admin:~/environment/sample-pipeline-artifacts $ git config --global user.email john@example.com
```

Next we can commit all newly added files to our local repository:

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git commit -a -m 'sample artifacts for OSS infrastructure pipeline'
```

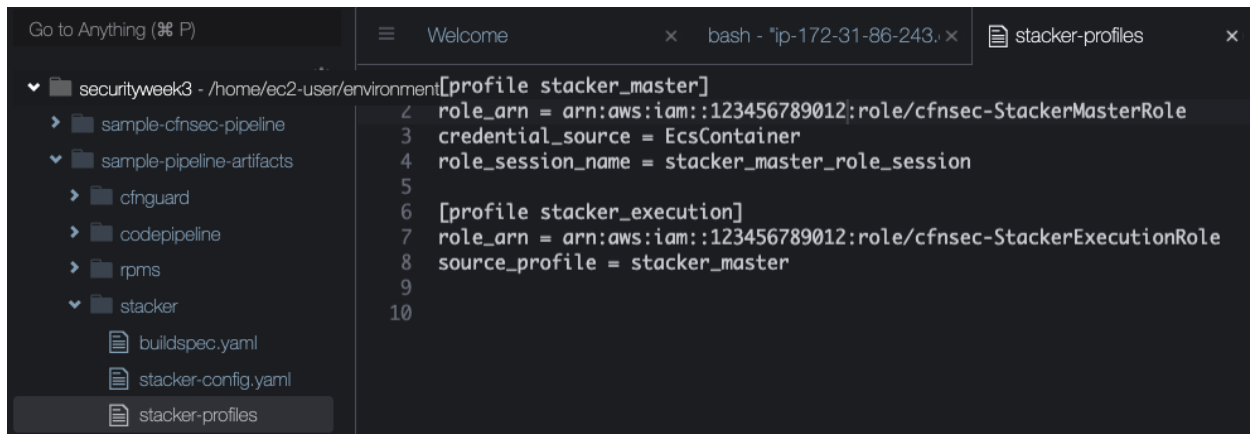
```
[master (root-commit) af7f464] sample artifacts for OSS infrastructure pipeline
```

```
...
```

```
15 files changed, 938 insertions(+)
create mode 100644 CODE_OF_CONDUCT.md
create mode 100644 CONTRIBUTING.md
create mode 100644 LICENSE
create mode 100644 README.md
create mode 100644 cfnguard/test.rules
create mode 100644 codepipeline/codepipeline-template.yaml
create mode 100644 codepipeline/stacker-execution-role-template.yaml
create mode 100644 rpms/cfn-guard-1.0-0.amzn2.x86_64.rpm
create mode 100644 rpms/cfn-guard-debuginfo-1.0-0.amzn2.x86_64.rpm
create mode 100644 stacker/buildspec.yaml
create mode 100644 stacker/stacker-config.yaml
create mode 100644 stacker/stacker-profiles
create mode 100644 templates/ec2-role-template.yaml
create mode 100644 templates/ec2-template.yaml
create mode 100644 templates/s3-bucket-template.yaml
```

5. Edit the Stacker-Profile file

Next, in the text editor in Cloud9 open the file *stacker/stacker-profiles* by double clicking it.



```
Go to Anything (% P) Welcome x bash - "ip-172-31-86-243." x stacker-profiles x
▼ securityweek3 - /home/ec2-user/environment
  ▼ sample-cfnsec-pipeline
  ▼ sample-pipeline-artifacts
    ▼ canguard
    ▼ codepipeline
    ▼ rpms
    ▼ stacker
      buildspec.yaml
      stacker-config.yaml
      stacker-profiles
[profile stacker_master]
2 role_arn = arn:aws:iam::123456789012:role/cfnsec-StackerMasterRole
3 credential_source = EcsContainer
4 role_session_name = stacker_master_role_session
5
6 [profile stacker_execution]
7 role_arn = arn:aws:iam::123456789012:role/cfnsec-StackerExecutionRole
8 source_profile = stacker_master
9
10
```

This file specifies the AWS profiles available in the CodeBuild environment running Stacker. Each profile refers to a specific **IAM Role**. The *stacker_master* profile is used to invoke Stacker from within the AWS CodeBuild environment. Stacker assumes that IAM role to access resources in the account where the pipeline is deployed. In order to perform stack deployments, Stacker needs one or more profiles. Since we're deploying to a single target account (which in our example is the same account where the pipeline is deployed), we define a single profile named *stacker_execution*. If we had multiple target accounts, we could define multiple profiles (e.g., *stacker_execution_accountA*, *stacker_execution_accountB*, etc.)

You'll need to modify the AWS account id **123456789012** specified in **stacker/stacker-profiles** and replace that value with your own AWS account ID.

To find your own aws account ID run the following command:

```
Admin:~/environment/sample-pipeline-artifacts (master) $ aws sts get-caller-identity | grep -i account | awk {'print $2'}
...
"546784563423",
```

This gives Stacker the target account for the deployment. After you update your account ID, save the edited file. Then run a git status to see that the file is recognized as modified by git:

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git status
...
On branch master
Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git restore <file>..." to discard changes in working directory)
    modified:   stacker/stacker-profiles
```

Next we can add the file to our staging area and commit it.

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git add .
```

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git commit -a -m 'Updated target AWS account ID in
stacker-profiles'
...
[master 9f2ec01] Updated target AWS account ID in stacker-profiles
1 file changed, 2 insertions(+), 2 deletions(-)
```

6. Edit the Stacker configuration file

In Cloud9, open the file ***stacker/stacker-config.yaml***. As discussed, this file contains parameters used by Stacker to create AWS resources for you.

Update parameter **s3_bucket_name** on **line 9** with a globally unique name **cfnsecsamples3bucketYour_AWS_ACCOUNT_ID** that will be created by Stacker as part of the CloudFormation stack deployment. This is just a simple private S3 bucket that will be created for you by Stacker. Also, note that the parameter *ec2_instance_type* indicates the EC2 instance type to deploy. By default, a “t2.micro” instance will be provisioned. Upon deployment, the EC2 instance will create and upload a simple file to the S3 bucket you specified. Once you have saved the file, make sure you commit it to the local repository.

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git status
...
On branch master
Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git restore <file>..." to discard changes in working directory)
    modified:   stacker/stacker-config.yaml
```

Add the file to the staging area and commit

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git add .
Admin:~/environment/sample-pipeline-artifacts (master) $ git commit -a -m 'updated s3 bucket name to be unique'
...
[master 939bf43] updated s3 bucket name to be unique
1 file changed, 1 insertion(+), 1 deletion(-)
```

7. Provision the pipeline

Now it's time to deploy the infrastructure pipeline using the sample code provided. Here we create two CloudFormation pipeline stacks as follows:

Pipeline stack #1: The first stack provisions the infrastructure pipeline we introduced earlier. The stack will provision a CodeCommit repo, a CodePipeline pipeline, and three CodeBuild projects.

The first CodeBuild project has a built-in [buildspec](#) configuration that will install and use CFN-Nag to check for security vulnerabilities in CloudFormation templates. The second CodeBuild project has a built-in buildspec configuration that will install and use CFN-Python-Lint to lint CloudFormation templates. The third CodeBuild project does not have a built-in buildspec configuration. Instead, a buildspec file will be pushed to the CodeCommit repo to instruct CodeBuild how it should install and run Stacker to provision and orchestrate CloudFormation stacks.

The buildspec file can either be built into the CloudFormation template or explicitly passed by developers. The advantage of building the buildspec file into the CloudFormation template is that it can help ensure that common standards are followed by teams who deploy code through a pipeline. (If you build the buildspec file into the CloudFormation template, the buildspec file cannot be edited or modified by developers.)

Before we go any further you will have to create an s3 bucket and upload the template. The s3 bucket should be in the following format mytemplates**YourAWSAccountID**:

```
Admin:~/environment/sample-pipeline-artifacts (master) $ aws s3 mb s3://mytemplates567893452376
...
make_bucket: mytemplates567893452376
```

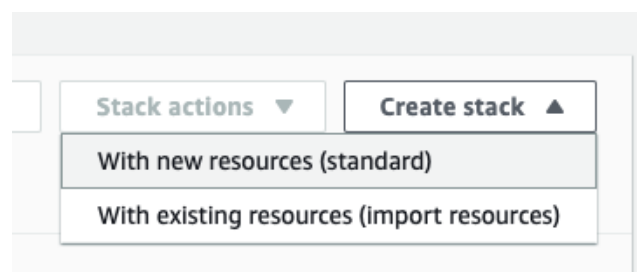
We then copy the codpipeline template codepipeline-template.yaml from the codepipeline folder into the s3 bucket

```
Admin:~/environment/sample-pipeline-artifacts (master) $ aws s3 cp codepipeline/codepipeline-template.yaml
s3://mytemplates567893452376/
```

8. Next we Provision the pipeline:

- Log in to your AWS Account
- Navigate to the [CloudFormation console](#).
- Create a stack named **cfnsec-codepipeline-stack** that represents the infrastructure pipeline

In the CloudFormation Stacks page, click the Create stack button with new resources:



On the following page (Create stack), complete the following sections:

- Prerequisite – Prepare template: leave the default radio button selected, **Template is ready**.
- Specify template – select the radio button, **Amazon S3 URL** a template file. Add your created bucket as your s3 url:

`https://mytemplatesYourAmazonAccountId.s3.us-east-1.amazonaws.com/codepipeline-template.yaml`

Create stack

Prerequisite - Prepare template

Prepare template

Every stack is based on a template. A template is a JSON or YAML file that contains configuration information about the AWS resources you want to include in the stack.

☒ Template is ready

☐ Use a sample template

☐ Create template in Designer

Specify template

A template is a JSON or YAML file that describes your stack's resources and properties.

Template source

Selecting a template generates an Amazon S3 URL where it will be stored.

☒ Amazon S3 URL

☐ Upload a template file

Amazon S3 URL

`https://mytemplates567893452376.s3.us-east-1.amazonaws.com/codepipeline-template.yaml`

Amazon S3 template URL

S3 URL: `https://mytemplates567893452376.s3.us-east-1.amazonaws.com/codepipeline-template.yaml`

[View in Designer](#)

Cancel

Next

- Click Next

On the page **Specify stack details**, complete the following sections:

- Stack name: enter the name **cfnsec-codepipeline-stack**.
- Parameters: leave the defaults or change as required (e.g., you can pick a name for you CodeCommit repository or accept the default “*samlerepo*”). Make sure to complete these two mandatory fields:
 - CodePipelineArtifactsS3BucketName: Specify a unique S3 bucket name that will be used by CodePipeline to store your build artifacts. We suggest you pick *cfnsec-pipeline-artifacts-XXXXXX* where X is a random number you choose or your AWS account id.
 - TargetAccount: Copy and paste your AWS account number in this field.

Stack name

Stack name

cfnsec-codepipeline-stack

Stack name can include letters (A-Z and a-z), numbers (0-9), and dashes (-).

Parameters

Parameters are defined in your template and allow you to input custom values when you create or update a stack.

CFNTemplatesPath
Relative path for CloudFormation templates

templates/

CodeCommitBranchName
name of codecommit branch to watch

master

CodeCommitRepoName
name of codecommit repo to create for infra code

samlerepo

CodePipelineArtifactsS3BucketName
Name of the CodePipeline S3 bucket to create to store temporary build artifacts

cfnsec-pipeline-artifacts-567893452376

Namespace
A prefix used to name provisioned resources

cfnsec

StackerConfigPath
Relative path for Stacker config files

stacker/

TargetAccount
ID of AWS Account where Stacker will provision resources

567893452376

Cancel Previous Next

- Click the Next button.

On the page **Configure stack options**, leave the defaults in place and select Next.

On the page **Review**, take a moment to review the inputs:

- Make sure that you have the unique S3 bucket name and your correct account number entered or this project will not work.

- Once you are satisfied that all details have been correctly entered, select the check box “**I acknowledge** that AWS CloudFormation might create IAM resources with custom names.” This is required because the template will create IAM roles for the various CodeBuild projects as well as the CodePipeline pipeline.
- Select Create stack. It will take a few minutes for your first stack to be created. You should see something like this:

The screenshot shows the AWS CloudFormation console for the stack 'cfnssec-codepipeline-stack'. The left sidebar shows a list of stacks, with 'cfnssec-codepipeline-stack' selected and showing a 'CREATE_IN_PROGRESS' status. The main panel shows the 'Events' tab for this stack, displaying a single event with the status 'CREATE_IN_PROGRESS' and the reason 'User Initiated'.

Timestamp	Logical ID	Status	Status reason
2020-10-17 10:10:52 UTC-0400	cfnssec-codepipeline-stack	CREATE_IN_PROGRESS	User Initiated

Once your stack has been created successfully, you will see output similar to:

The screenshot shows the AWS CloudFormation console for the stack 'cfnssec-codepipeline-stack'. The left sidebar shows the stack status as 'CREATE_COMPLETE'. The main panel shows the 'Events' tab, displaying 65 events. The first event shown is 'CodeBuildLogsPolicy' with a status of 'CREATE_COMPLETE'.

Timestamp	Logical ID	Status	Status reason
2020-10-17 10:11:26 UTC-0400	CodeBuildLogsPolicy	CREATE_COMPLETE	-

- Before you create the second stack, navigate over to the [CodeCommit Console](#). Select **Repositories** and you should see a sample repository created, like this (the name of the repository will be different if you have changed the default value):

The screenshot shows the AWS CodeCommit console 'Repositories' page. It displays a table with one repository named 'samplerepo', which is a 'CodeCommit repository' created '3 minutes ago'. The 'Clone URL' column shows links for HTTPS, SSH, and HTTPS (GRC).

Name	Description	Last modified	Clone URL
○ samplerepo	CodeCommit repository	3 minutes ago	HTTPS SSH HTTPS (GRC)

Next, navigate to the [IAM console](#) and select **Roles**. You'll be able to see the five roles that your stack created:

Q cfnsec			
	Role name ▼	Trusted entities	Last activity ▼
<input type="checkbox"/>	cfnsec-CodeBuildCFNGuardRole	AWS service: codebuild	None
<input type="checkbox"/>	cfnsec-CodeBuildCFNLintRole	AWS service: codebuild	None
<input type="checkbox"/>	cfnsec-CodeBuildCFNNagRole	AWS service: codebuild	None
<input type="checkbox"/>	cfnsec-CodeBuildDeployerRole	AWS service: codebuild	None
<input type="checkbox"/>	cfnsec-CodePipelineRole	AWS service: codepipeline	None
<input type="checkbox"/>	cfnsec-StackerMasterRole	Account: 447333226489	None

Let's look more closely at each of the roles and make sure that we understand the purpose of each role:

- **cfnsec-CodeBuildCFNLintRole:** This is an IAM service role created for the CFN-Python-Lint CodeBuild environment
- **cfnsec-CodeBuildCFNNagRole:** This is an IAM service role created for the CFN-Nag CodeBuild environment
- **cfnsec-CodeBuildCFNGuardRole:** This is an IAM service role created for the CFN-Guard CodeBuild environment
- **cfnsec-CodeBuildDeployerRole:** This is an IAM service role created for the Stacker CodeBuild environment
- **cfnsec-CodePipelineRole:** This is an IAM service role for CodePipeline to allow the pipeline to perform tasks such as read/write artifacts from/to the artifacts S3 bucket, to trigger the various CodeBuild environments
- **cfnsec-StackerMasterRole:** This is an IAM role used to launch Stacker, which allows Stacker to assume the cfnsec-StackerExecutionRole (to be created next) to deploy AWS resources via CloudFormation on the various target accounts. For simplicity, in our example we are using a single account, i.e., the same account where the pipeline resides.

The file **stacker/buildspec.yaml** contains the AWS CodeBuild buildspec to install and invoke Stacker for CloudFormation template provisioning. The **buildspec.yaml** file defines the steps taken by AWS CodeBuild to test the AWS CloudFormation template prior to deploying it into a production environment. As mentioned before, there are two main ways in which you can specify your CodeBuild buildspec configuration. The first is to specify a buildspec inside the CloudFormation template (built-in buildspec) as we did for CFN-Nag and CFN-Python-Lint.

The code snippet below shows the CodeBuild project created via CloudFormation **codepipeline-template.yaml** for the CFN-Python-Lint CodePipeline action. Notice property **BuildSpec** below and

how we used it to install (see `pip install cfn-lint`) and invoke the CFN-Python-Lint tool from within our CodeBuild Python 3.6.5 environment.

```
CFNLintCodeBuildProject:
  Type: AWS::CodeBuild::Project
  Properties:
    Name: !Sub ${Namespace}-cfn-lint-code-build-project
    Description: CodeBuild Project to validate CloudFormation templates using cnf-python-lint
    Artifacts:
      Type: CODEPIPELINE
    Environment:
      Type: LINUX_CONTAINER
      ComputeType: BUILD_GENERAL1_SMALL
      Image: aws/codebuild/python:3.6.5
      EnvironmentVariables:
        - Name: CFNTemplatesPath
          Value: !Ref CFNTemplatesPath
    ServiceRole:
      !GetAtt CFNNagCodeBuildServiceRole.Arn
    Source:
      Type: CODEPIPELINE
      BuildSpec: |
        version: 0.2
        phases:
          install:
            commands:
              - pip install --upgrade pip
              - env && ls -l && python --version
              - pip install cfn-lint
              - cfn-lint ${CFNTemplatesPath}*.yaml
```

```
CFNPNagCodeBuildProject:
  Type: AWS::CodeBuild::Project
  Properties:
    Name: !Sub ${Namespace}-cfn-nag-code-build-project
    Description: CodeBuild Project to validate CloudFormation templates using CFN-Nag
    Artifacts:
      Type: CODEPIPELINE
    Environment:
      Type: LINUX_CONTAINER
      ComputeType: BUILD_GENERAL1_SMALL
      # With the image below we must specify a runtime-version in the Builds spec (see below)
      Image: aws/codebuild/amazonlinux2-x86_64-standard:1.0
      EnvironmentVariables:
        - Name: CFNTemplatesPath
          Value: !Ref CFNTemplatesPath
    ServiceRole:
      !GetAtt CFNLintCodeBuildServiceRole.Arn
    Source:
      Type: CODEPIPELINE
      BuildSpec: |
        version: 0.2
        phases:
          install:
            runtime-versions:
```

```

    ruby: 2.6
  commands:
    - env && ls -l && ruby -v
    - gem install cfn-nag
    - cfn_nag_scan -v
    - cfn_nag_scan --input-path $CFNTemplatesPath

CFNGuardCodeBuildProject:
  Type: AWS::CodeBuild::Project
  Properties:
    Name: !Sub ${Namespace}-cfn-guard-code-build-project
    Description: CodeBuild Project to validate CloudFormation templates using cnf-guard
    Artifacts:
      Type: CODEPIPELINE
    Environment:
      Type: LINUX_CONTAINER
      ComputeType: BUILD_GENERAL1_SMALL
      # With the image below we must specify a runtime-version in the Builds spec (see below)
      Image: aws/codebuild/amazonlinux2-x86_64-standard:1.0
      EnvironmentVariables:
        - Name: CFNTemplatesPath
          Value: !Ref CFNTemplatesPath
    ServiceRole:
      !GetAtt CFNGuardCodeBuildServiceRole.Arn
    Source:
      Type: CODEPIPELINE
      BuildSpec: |
        version: 0.2
        phases:
          install:
            runtime-versions:
              ruby: 2.6
            commands:
              - rm -rf /var/cache/yum
              - yum clean all && yum makecache fast
              - yum -y install rust cargo
              - rpm -ivh $CFNTemplatesPath/cfn-guard-1.0-0.amzn2.x86_64.rpm
              - which cfn-guard
              - cfn-guard check --rule_set $CFNTemplatesPath/test.rules --template ${CFNTemplatesPath}/ec2-role-template.yaml
              - cfn-guard check --rule_set $CFNTemplatesPath/test.rules --template ${CFNTemplatesPath}/ec2-template.yaml
              - cfn-guard check --rule_set $CFNTemplatesPath/test.rules --template ${CFNTemplatesPath}/s3-bucket-template.yaml

```

Using the **BuildSpec** property in CloudFormation means that developers pushing code to the pipeline cannot modify the behavior of the CodeBuild environment and how we use CFN-Python-Lint to lint CloudFormation templates.

Another way to configure a CodeBuild environment is to allow developers to create a custom **buildspec.yml** file and push it to the CodeCommit repository. This gives them more flexibility and control over the CodeBuild execution environment so they, for instance, install and

run software as needed. In cases where different teams own the pipeline vs. the pipeline artifacts the pipeline builder team should plan carefully what level of customization they want to offer developers vs. what build steps are immutable.

Below is the custom **buildspec.yml** YAML file we use for the CodeBuild project that runs Stacker. Notice that we define some environment variables and run a few commands to configure AWS credentials and to install (**pip install stacker==1.7.0**) and run Stacker (**stacker build**).

```
version: 0.2

env:
  variables:
    stacker_master_profile_name: "stacker_master"
    stacker_profiles_file: "stacker-profiles"
    stacker_orchestration_file: "stacker-config.yaml"

phases:

  pre_build:
    commands:
      - pip install --upgrade pip
      - pip install stacker==1.7.0
      - env && ls -lha && python --version

  build:
    commands:
      - export AWS_CONFIG_FILE="${CODEBUILD_SRC_DIR}/${StackerConfigPath}/${stacker_profiles_file}"
      - echo "AWS_CONFIG_FILE=${AWS_CONFIG_FILE}"
      - stacker build "${CODEBUILD_SRC_DIR}/${StackerConfigPath}/${stacker_orchestration_file}" --profile
        $stacker_master_profile_name --recreate-failed
```

Pipeline stack #2: The second stack we're going to create provisions the IAM role required by Stacker to deploy AWS resources in a given target AWS account. We use separate templates for the pipeline and the Stacker execution role, because the latter can now be deployed to multiple target accounts if we wish the pipeline to provision stacks into those accounts.

Since we're using a single AWS account in this post, we will create the Stacker execution role in the same account as the pipeline. The IAM policies in the role must set permissions to deploy all types of AWS resources referenced in the template. In our case, Stacker needs permissions to create an S3 bucket, an EC2 instance, IAM roles, etc. Please check the template for further details.

10. Create your second stack the same way you created the first one, named **cfnsec-stacker-execution-role-stack**, using the file **codepipeline/stacker-execution-role-template.yaml**. This second stack will create an IAM role in your account so that Stacker can deploy outputs in the same account where it lives. As stated above, while we are using a single account for this project, this approach will also work for multiple accounts.

Upload the file to your s3 bucket:

```
Admin:~/environment/sample-pipeline-artifacts (master) $ aws s3 cp codepipeline/stacker-execution-role-template.yaml s3://mytemplates567893452376/
```

On the **CloudFormation** Stacks page, select Create stack.

- On the following page (Create stack), complete the following sections:
- Prerequisite – Prepare template: leave the default radio button selected, **Template is ready**.
- Specify template – select the radio button, **Amazon S3 URL** a template file. Add your created bucket as your s3 url:

`https://mytemplatesYourAmazonAccountId.s3.us-east-1.amazonaws.com/stacker-execution-role-template.yaml`

- Click Next.

Prerequisite - Prepare template

Prepare template
Every stack is based on a template. A template is a JSON or YAML file that contains configuration information about the AWS resources you want to include in the stack.

☒ Template is ready

☐ Use a sample template

☐ Create template in Designer

Specify template

A template is a JSON or YAML file that describes your stack's resources and properties.

Template source
Selecting a template generates an Amazon S3 URL where it will be stored.

☒ Amazon S3 URL

☐ Upload a template file

Amazon S3 URL

`https://mytemplates567893452376.s3.us-east-1.amazonaws.com/stacker-execution-role-template.yaml`

Amazon S3 template URL

S3 URL: `https://mytemplates567893452376.s3.us-east-1.amazonaws.com/stacker-execution-role-template.yaml` [View in Designer](#)

[Cancel](#) [Next](#)

On the page **Specify stack details**, complete the following sections:

- Stack name: enter the name **cfnsec-stacker-execution-role-stack**
- Parameters: leave the defaults. However, make sure to complete these two fields:
 - Namespace: leave the default value, cfnsec, as is.
 - StackMasterAccountID: Copy and paste your AWS account number in this field.

Specify stack details

Stack name

Stack name

cfnsec-stacker-execution-role-stack

Stack name can include letters (A-Z and a-z), numbers (0-9), and dashes (-).

Parameters

Parameters are defined in your template and allow you to input custom values when you create or update a stack.

Namespace

A prefix used to name provisioned resources

cfnsec

StackerMasterAccountid

Account ID where the StackerMasterRole lives

567893452376

Cancel Previous Next

- Select Next

On the **Configure stack options** page, leave the defaults in place and select Next

On the **Review page**, take a moment to review the inputs.

- Make sure that you have the unique S3 bucket name and your correct account number entered or this project will not work.
- At the bottom of the page, click the tick box **"I acknowledge that AWS CloudFormation might create IAM resources with custom names."**
- Click Create Stack. It will take a minute or two for your stack to be created. You should see something like this:

Stacks (4)

Filter by stack name

Active View nested

cfnsec-stacker-execution-role-stack
2020-10-17 10:27:12 UTC-0400
CREATE_IN_PROGRESS

cfnsec-codepipeline-stack
2020-10-17 10:10:52 UTC-0400
CREATE_COMPLETE

aws-cloud9-securityweek3-fe46d61730304150b110
b60e76490615

cfnsec-stacker-execution-role-stack

Delete Update Stack actions Create stack

Stack info Events Resources Outputs Parameters Template Change sets

Events (1)

Search events

Timestamp	Logical ID	Status	Status reason
2020-10-17 10:27:12 UTC-0400	cfnsec-stacker-execution-role-stack	CREATE_IN_PROGRESS	User initiated

When your stack has been successfully built, the console screen will look like this:

Stacks (4)

Filter by stack name

Active View nested

cfnsec-stacker-execution-role-stack
2020-10-17 10:27:12 UTC-0400
CREATE_COMPLETE

cfnsec-codepipeline-stack
2020-10-17 10:10:52 UTC-0400
CREATE_COMPLETE

aws-cloud9-securityweek3-fe46d61730304150b110
b60e76490615

Events (8)

Search events

Timestamp	Logical ID	Status	Status reason
2020-10-17 10:27:46 UTC-0400	cfnsec-stacker-execution-role-stack	CREATE_COMPLETE	-


11. Using the AWS CodePipeline infrastructure pipeline

Now that the pipeline is ready and we have configured the pipeline artifacts, it's time to use the pipeline to deploy our stacks.

Prerequisites

Next we will setup our local repository sample-pipeline-artifacts to upload all the files to our sample repository. From within your AWS Account, navigate to the CodeCommit console and look for the repo named samplerepo.

- Under the heading Clone URL, click the blue text HTTPS. This action will copy the HTTPS address to your clipboard.



At this point, you should have a local git repository with the **stacker/** and **templates/** folders (created earlier). Add a Git remote to your CodeCommit repository like this:

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git remote add origin https://git-codecommit.us-east-1.amazonaws.com/v1/repos/samplerepo
```

Next we need to add two more files to our templates folder so that we can run cfn-guard

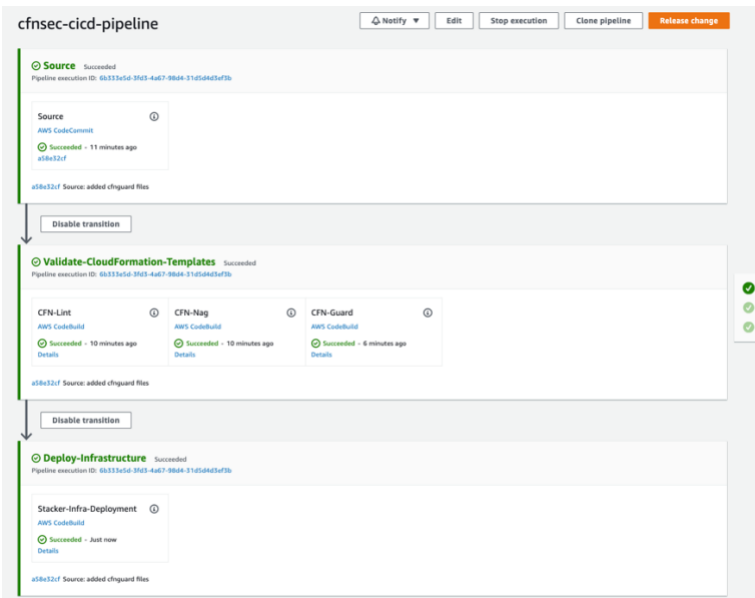
```
Admin:~/environment/sample-pipeline-artifacts (master) $ cp cfn-guard/test.rules templates/
Admin:~/environment/sample-pipeline-artifacts (master) $ cp rpms/cfn-guard-1.0-0.amzn2.x86_64.rpm templates/
Admin:~/environment/sample-pipeline-artifacts (master) $ git add .
Admin:~/environment/sample-pipeline-artifacts (master) $ git status
...
On branch master
Changes to be committed:
  (use "git restore --staged <file>..." to unstage)
    new file:   templates/cfn-guard-1.0-0.amzn2.x86_64.rpm
    new file:   templates/test.rules

Admin:~/environment/sample-pipeline-artifacts (master) $ git commit -a -m 'added cfn-guard files'
[master a58e32c] added cfn-guard files
2 files changed, 9 insertions(+)
create mode 100644 templates/cfn-guard-1.0-0.amzn2.x86_64.rpm
create mode 100644 templates/test.rules
```

Make sure your local Git repo does not have any pending changes to be committed (git status). Push your code to the CodeCommit repository's master branch:

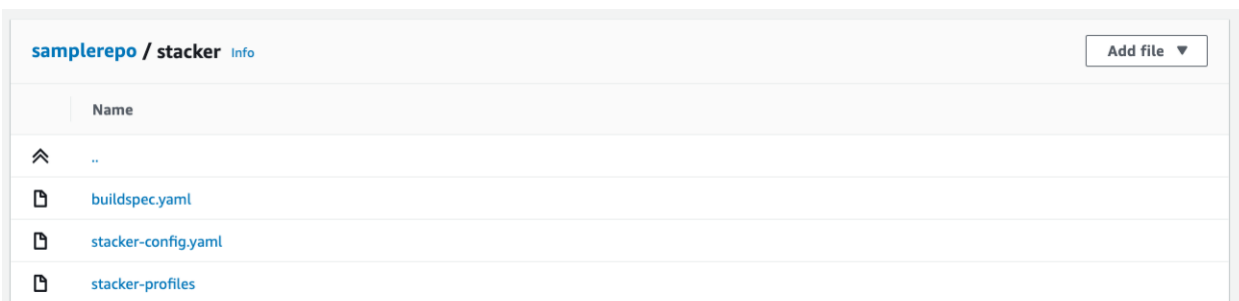
```
Admin:~/environment/sample-pipeline-artifacts (master) $ git push origin master:master
```

After pushing your code, navigate to the AWS CodePipeline console to verify that the pipeline was triggered successfully. Once in the CodePipeline console, select **Pipelines** then your pipeline **cfnsec-cicd-pipeline**. If all goes well, your templates will be validated and deployed and the pipeline should look like the screenshot below:



12. So what have we deployed exactly? Well, the three stacks we mentioned previously: an S3 bucket, an IAM role, and an EC2 instance (plus other resources, e.g., security group, EC2 profile) that assumes the IAM role and writes a file to the S3 bucket. You should have a file named 'hello.txt' in the S3 bucket you specified in **stacker/stacker-config.yaml**.

Have a look now at your sample repo in the AWS CodeCommit Console, to review what you pushed through your serverless pipeline:



Earlier in this post we introduced the file structure and explained the files that will be pushed out to the codepipeline and their purpose. Let's briefly recap:

- **buildspec.yaml**: This is the AWS CodeBuild buildspec. As discussed, its purpose is to install and invoke Stacker for CloudFormation resource provisioning
- **stacker-config.yaml**: This is the stacker config file containing stack descriptions and input parameters
- **stacker-profiles**: This file contains the AWS profiles that Stacker used to assume roles during deployment

13. Testing pipeline failure scenarios

Let's make a change to one of the templates under `templates/` to cause the pipeline template verification stage to fail. For example, specify an invalid type for the `S3BucketName` input parameter in the template `s3-bucket-template.yaml`, like this:

```
AWSTemplateFormatVersion: "2010-09-09"
Description: Creates an encrypted non-public S3 bucket to store CI/CD artifacts

Parameters:

  S3BucketName:
    Type: *InvalidType*
    Description: Unique name to assign to S3 bucket

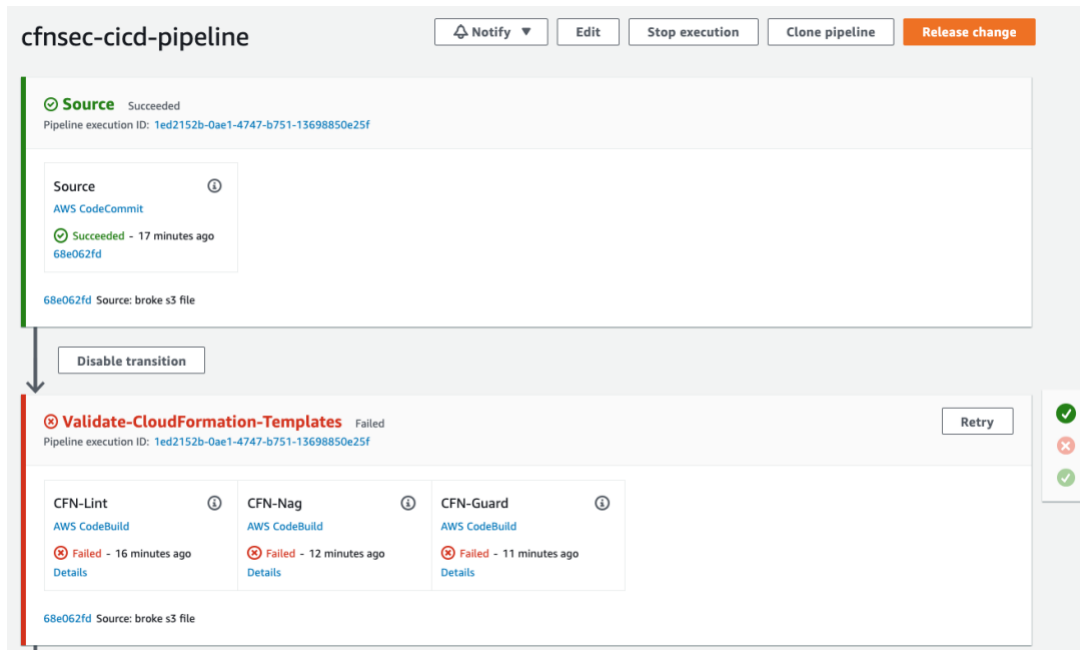
Resources:

  Bucket:
    Type: AWS::S3::Bucket
    DeletionPolicy: Delete
```

Commit and push your code to the CodeCommit repository again which will trigger the pipeline.

```
Admin:~/environment/sample-pipeline-artifacts (master) $ git add .
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git commit s3-bucket-template.yaml -m 'broke s3 file'
...
[master 68e062f] broke s3 file
1 file changed, 1 insertion(+), 1 deletion(-)
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git push origin master:master
Enumerating objects: 7, done.
Counting objects: 100% (7/7), done.
Delta compression using up to 2 threads
Compressing objects: 100% (4/4), done.
Writing objects: 100% (4/4), 363 bytes | 363.00 KiB/s, done.
Total 4 (delta 3), reused 0 (delta 0)
To https://git-codecommit.us-east-1.amazonaws.com/v1/repos/samplerepo
a58e32c..68e062f master -> master
```

Wait for the pipeline to fail as CFN-Python-Lint checks your template against CloudFormation specs, as in the screenshot below:



For details on why the pipeline failed, click on the **Details** link in the CodeBuild action for CFN- -Lint. You should see a screen similar to the one below.

```
120
121 [Container] 2020/10/16 23:28:53 Running command cfn-lint ${CFNTemplatesPath}*.yaml
122 E0000 did not find expected alphabetic or numeric character
123 templates/s3-bucket-template.yaml:7:23
124
125
126 [Container] 2020/10/16 23:28:55 Command did not exit successfully cfn-lint ${CFNTemplatesPath}*.yaml exit status 2
127
```

That's exactly what we wanted: CFN-Lint is helping us make sure that our templates are valid before letting the pipeline deploy them.

Let's revert the last commit again. Go back to your s3 file and replace the `*invalidType*` with `String`. Then run the commands below which will trigger the pipeline again and should provide a successful build.

```
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git add s3-bucket-template.yaml
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git commit s3-bucket-template.yaml -m 'reverted breaking the s3 bucket'
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git push origin master:master
...
Enumerating objects: 7, done.
Counting objects: 100% (7/7), done.
Delta compression using up to 2 threads
Compressing objects: 100% (4/4), done.
Writing objects: 100% (4/4), 367 bytes | 367.00 KiB/s, done.
Total 4 (delta 3), reused 0 (delta 0)
To https://git-codecommit.us-east-1.amazonaws.com/v1/repos/samplerepo
47a3514..8a88bd4 master -> master
```

14. CFN-guard build failure

We will now look at what `cnf-guard`. Below is the `test.rules` files from the `cfnguard` folder:

```
AWS::EC2::Instance IamInstanceProfile == ECInstanceProfile
AWS::EC2::Instance ImageId == LatestAmiId
AWS::EC2::Instance InstanceType == /.nano|.micro|.small/
AWS::EC2::Instance SecurityGroups == ["InstanceSecurityGroup"]
AWS::EC2::Instance UserData == { "Fn::Base64": "#!/bin/bash -xe\nexport INSTANCE_ID=$(curl -s http://169.254.169.254/latest/meta-data/instance-id)\necho \"Hello from EC2 instance $INSTANCE_ID\" > hello.txt\naws s3 cp hello.txt s3://${S3BucketName}\n" }
AWS::EC2::SecurityGroup GroupDescription == Enable SSH access via port 22
AWS::EC2::SecurityGroup SecurityGroupIngress == [{"CidrIp": "10.15.45.31/32", "FromPort": 22, "IpProtocol": "tcp", "ToPort": 22}]
AWS::IAM::InstanceProfile Path == /
AWS::IAM::InstanceProfile Roles == ["EC2IamRoleName"]
```

This creates guardrails which define that an `InstanceType` can only be `nano`, `micro` or `small`. You can define for example what your AMI ID should be as well as define specific security groups and much more. For example if you have a security standard defined, you can model the `cnf-guard` rules after it. If a developer chooses to use a larger instance then defined, `cnf-guard` will fail the build.

Let run through the last example of this lab.

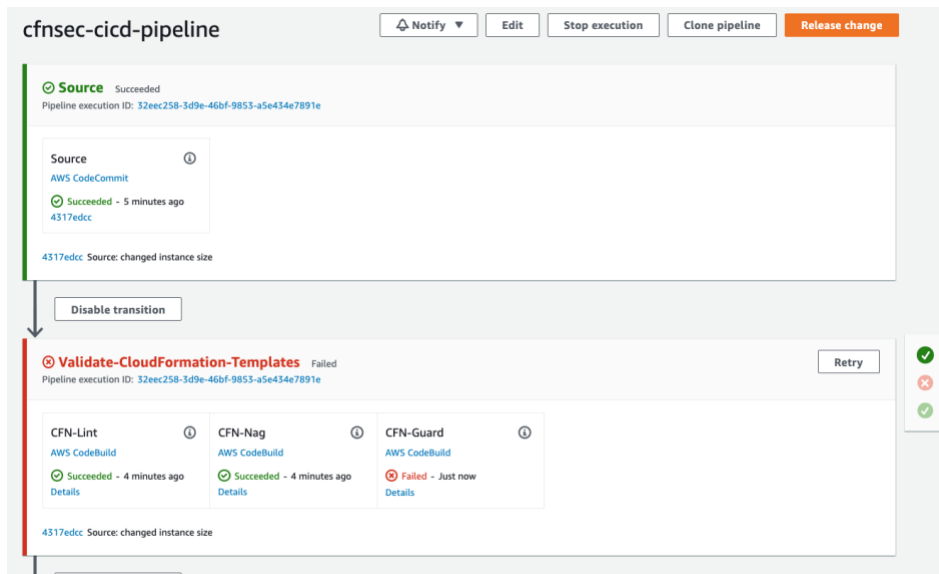
We will alter the following file `ec2-template.yml` which is in the `templates` folder. We will change the `InstanceType` to `m5.2xlarge`.

```
EC2Instance:
  Type: AWS::EC2::Instance
  Properties:
    InstanceType: m5.2xlarge
    SecurityGroups: [!Ref 'InstanceSecurityGroup']
    ImageId: !Ref 'LatestAmiId'
    IamInstanceProfile: !Ref ECInstanceProfile
```

Add the file to your staging area, commit and push the file to the repository.

```
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git add .
Admin:~/environment/sample-pipeline-artifacts/templates (master) $ git commit ec2-template.yaml -m 'changed instance size'
...
[master 4317edc] changed instance size
 1 file changed, 1 insertion(+), 1 deletion(-)
git push origin master:master
Enumerating objects: 7, done.
Counting objects: 100% (7/7), done.
Delta compression using up to 2 threads
Compressing objects: 100% (4/4), done.
Writing objects: 100% (4/4), 370 bytes | 370.00 KiB/s, done.
Total 4 (delta 3), reused 0 (delta 0)
To https://git-codecommit.us-east-1.amazonaws.com/v1/repos/samplerepo
 8a88bd4..4317edc master -> master
```

This will trigger another build. You will see that this test is passing cfn-lint as well as cfn-nag as our template is syntactically correct and we are not violating the rules of cfn-nag but we will fail cfn-guard as it does not meet our specifications.



If you investigate the build you will see that cfn-guard caught our change in terms of a not permitted instance size.

```
[Container] 2020/10/17 16:25:07 Running command cfn-guard check --rule_set ${CFNTemplatesPath}/test.rules --template ${CFNTemplatesPath}/ec2-template.yaml
[EC2Instance] failed because [InstanceType] is [m5.2xlarge] and the permitted pattern is [.nano|.micro|.small]
Number of failures: 1

[Container] 2020/10/17 16:25:07 Command did not exit successfully cfn-guard check --rule_set ${CFNTemplatesPath}/test.rules --template ${CFNTemplatesPath}/ec2-template.yaml exit status 2
```

Conclusion

In this lab, we showed you how to leverage popular open source tools in conjunction with AWS Code services such as CodePipeline, CodeCommit, and CodeBuild to build, validate and deploy arbitrary infrastructure stacks. We used CFN-NAG, CFN-Lint and CFN-guard to validate and guard the templates, and Stacker to perform the deployment of the CloudFormation stacks.